

DATA698 - Data Cleaning and Look Through

Max Wagner

```
#libs
require("ROSE")

## Warning: package 'ROSE' was built under R version 3.3.3

require("pROC")
require("rpart")
require("rpart.plot")

## Warning: package 'rpart.plot' was built under R version 3.3.3

require("caret")
require("randomForest")

## Warning: package 'randomForest' was built under R version 3.3.3

require("e1071")

#main file
cc <- data.frame(read.csv("data/cc.csv"))
cc <- cc[,c(2:31)]

#check balance
fraud <- nrow(cc[cc$Class == 1,])
notFraud <- nrow(cc) - fraud
paste("fraud: ", fraud, "|| not fraud: ", notFraud)

## [1] "fraud: 492 || not fraud: 284315"

#make the size smaller for easier use
set.seed(65)
cc_simp <- cc[sample(nrow(cc), 25000), ]
fraud <- nrow(cc_simp[cc_simp$Class == 1,])
notFraud <- nrow(cc_simp) - fraud
paste("fraud: ", fraud, "|| not fraud: ", notFraud)

## [1] "fraud: 43 || not fraud: 24957"

#split data for testing models
trainLength <- floor(.7*nrow(cc_simp))
testLength <- nrow(cc_simp) - trainLength

train_model <- cc_simp[1:trainLength,]
train_eval <- cc_simp[(trainLength + 1):nrow(cc_simp),]

#functions for sd and se
mysd <- function(predict, target) {
  diff_sq <- (predict - mean(target))^2
  return(mean(sqrt(diff_sq)))
}
```

```

}

myse <- function(predict, target) {
  diff_sq <- (predict - target)^2
  return(mean(sqrt(diff_sq)))
}

#Model1 - Multiple Linear Regression - Base Line
mlr1 <- glm(Class~., data = train_model)
BIC(mlr1)

## [1] -71401.37

predict_mlr1 <- predict(mlr1, train_eval, type = 'response')
table(train_eval$Class, predict_mlr1 > 0.5)

##
##      FALSE TRUE
##    0  7485    1
##    1   10    4

mysd(predict_mlr1, train_eval$Class)


## [1] 0.00414942

myse(predict_mlr1, train_eval$Class)

## [1] 0.004008179

auc_mlr1 <- roc(train_eval$Class, predict_mlr1)
plot(auc_mlr1)

```



```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_mlr1)
##
## Data: predict_mlr1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9949
#Model2 - Poisson Model
poisson1 <- glm(Class ~ ., family = "poisson", data = train_model)


## Warning: glm.fit: fitted rates numerically 0 occurred
BIC(poisson1)

## [1] 487.1279
predict_poisson1 <- predict(poisson1, train_eval, type = 'response')
table(train_eval$Class, predict_poisson1 > 0.5)

##
##      FALSE TRUE
## 0    7484     2
## 1      12     2
mysd(predict_poisson1, train_eval$Class)

## [1] 0.002839854
myse(predict_poisson1, train_eval$Class)
```

```
## [1] 0.002619012
auc_poisson <- roc(train_eval$Class, predict_poisson1)
plot(auc_poisson)
```



```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_poisson1)
##
## Data: predict_poisson1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9771
```

```
#logit model
logit1 <- glm(Class ~., family = binomial(link='logit'), data = train_model)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


```
BIC(logit1)
```

```
## [1] 440.7449
```

```
predict_logit1 <- predict(logit1, train_eval, type = 'response')
table(train_eval$Class, predict_logit1 > 0.5)
```

```
##
##      FALSE TRUE
## 0  7484     2
## 1     7     7
```

```
auc_logit1 <- roc(train_eval$Class, predict_logit1)
plot(auc_logit1)
```



```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_logit1)
##
## Data: predict_logit1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9859
# backward stepwise
stepwise1 <- glm(Class ~ ., data = train_model)
backward <- step(stepwise1, trace = 0)
BIC(backward)

## [1] -71455.25

predict_backward <- predict(backward, train_eval, type = 'response')
table(train_eval$Class, predict_backward > 0.5)

##
##      FALSE TRUE
## 0   7485    1
## 1     10    4

mysd(predict_backward, train_eval$Class)


## [1] 0.0041126
```

```
myse(predict_backward, train_eval$Class)
```

```
## [1] 0.003986638
```

```
auc_backward <- roc(train_eval$Class, predict_backward)
```

```
plot(auc_backward)
```



```
##
```

```
## Call:
```

```
## roc.default(response = train_eval$Class, predictor = predict_backward)
```

```
##
```

```
## Data: predict_backward in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
```

```
## Area under the curve: 0.9951
```

```
#forward stepwise
```

```
stepwise2 <- glm(Class ~ 1, data = train_model)
```

```
forward <- step(stepwise2, scope = list(lower=formula(stepwise2), upper=formula(stepwise1)), direction = "both",  
BIC(forward))
```

```
## [1] -71455.25
```

```
predict_forward <- predict(forward, train_eval, type = 'response')
```

```
table(train_eval$Class, predict_forward > 0.5)
```

```
##
```

```
##      FALSE TRUE
```

```
## 0  7485    1
```

```
## 1    10    4
```


```
mysd(predict_forward, train_eval$Class)

## [1] 0.0041126
myse(predict_forward, train_eval$Class)

## [1] 0.003986638
auc_forward <- roc(train_eval$Class, predict_forward)
plot(auc_forward)

##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_forward)
##
## Data: predict_forward in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9951

# decision tree
decision <- rpart(Class ~ ., data = train_model, method = "class")
prp(decision)
```



```
predict_decision <- predict(decision, train_eval, type = "class")
confusionMatrix(train_eval$Class, predict_decision)
```

```
## Confusion Matrix and Statistics
##
##           Reference
```

```

## Prediction      0      1
##              0 7484    2
##              1      8    6
##
##              Accuracy : 0.9987
##              95% CI : (0.9975, 0.9994)
##      No Information Rate : 0.9989
##      P-Value [Acc > NIR] : 0.8160
##
##              Kappa : 0.5448
##      McNemar's Test P-Value : 0.1138
##
##              Sensitivity : 0.9989
##              Specificity : 0.7500
##      Pos Pred Value : 0.9997
##      Neg Pred Value : 0.4286
##              Prevalence : 0.9989
##      Detection Rate : 0.9979
##      Detection Prevalence : 0.9981
##      Balanced Accuracy : 0.8745
##
##      'Positive' Class : 0
##
#decision tree random forest (kind of broke with raw data)
rforest <- randomForest(Class ~ ., data = train_model)


## Warning in randomForest.default(m, y, ...): The response has five or fewer
## unique values. Are you sure you want to do regression?

predict_rforest <- predict(rforest, train_eval)
table(train_eval$Class, predict_rforest > 0.5)

##
##      FALSE TRUE
##      0 7485    1
##      1      6    8

auc_rforest <- roc(train_eval$Class, predict_rforest)
plot(auc_rforest)

```





```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_rforest)
##
## Data: predict_rforest in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9712

#svm (kind of broken with raw data)
svm <- svm(Class ~ ., data = train_model)
predict_svm <- predict(svm, train_eval)
table(train_eval$Class, predict_svm > 0.5)

##
##      FALSE
##  0  7486
##  1    14

auc_svm <- roc(train_eval$Class, predict_svm)
plot(auc_svm)
```



```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_svm)
##
## Data: predict_svm in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9505

#tables only
table(train_eval$Class, predict_mlr1 > 0.5)

##
##      FALSE TRUE
## 0  7485     1
## 1    10     4

table(train_eval$Class, predict_poisson1 > 0.5)

##
##      FALSE TRUE
## 0  7484     2
## 1    12     2

table(train_eval$Class, predict_logit1 > 0.5)

##
##      FALSE TRUE
## 0  7484     2
## 1     7     7
```

```
table(train_eval$Class, predict_backward > 0.5)
```

```
##
##      FALSE TRUE
##  0  7485     1
##  1     10     4
```

```
table(train_eval$Class, predict_forward > 0.5)
```

```
##
##      FALSE TRUE
##  0  7485     1
##  1     10     4
```

```
confusionMatrix(train_eval$Class, predict_decision)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0     1
##              0 7484     2
##              1     8     6
##
##              Accuracy : 0.9987
##              95% CI : (0.9975, 0.9994)
##              No Information Rate : 0.9989
##              P-Value [Acc > NIR] : 0.8160
##
##              Kappa : 0.5448
##              Mcnemar's Test P-Value : 0.1138
##
##              Sensitivity : 0.9989
##              Specificity : 0.7500
##              Pos Pred Value : 0.9997
##              Neg Pred Value : 0.4286
##              Prevalence : 0.9989
##              Detection Rate : 0.9979
##              Detection Prevalence : 0.9981
##              Balanced Accuracy : 0.8745
##
##              'Positive' Class : 0
##
```


```
table(train_eval$Class, predict_rforest > 0.5)
```

```
##
##      FALSE TRUE
##  0  7485     1
##  1     6     8
```


```
table(train_eval$Class, predict_svm > 0.5)
```

```
##
##      FALSE
##  0  7486
##  1    14
```


```
#AUC plots only  
plot(auc_mlr1)
```




```
##  
## Call:  
## roc.default(response = train_eval$Class, predictor = predict_mlr1)  
##  
## Data: predict_mlr1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).  
## Area under the curve: 0.9949  
plot(auc_poisson)
```



```
##  
## Call:  
## roc.default(response = train_eval$Class, predictor = predict_poisson1)  
##  
## Data: predict_poisson1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).  
## Area under the curve: 0.9771  
plot(auc_logit1)
```




```
##  
## Call:  
## roc.default(response = train_eval$Class, predictor = predict_logit1)  
##  
## Data: predict_logit1 in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).  
## Area under the curve: 0.9859  
plot(auc_backward)
```




```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_backward)
##
## Data: predict_backward in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9951
plot(auc_forward)


##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_forward)
##
## Data: predict_forward in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9951
prp(decision)
```



```
plot(auc_rforest)
```

```
##  
## Call:  
## roc.default(response = train_eval$Class, predictor = predict_rforest)  
##  
## Data: predict_rforest in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).  
## Area under the curve: 0.9712  
plot(auc_svm)
```



```
##
## Call:
## roc.default(response = train_eval$Class, predictor = predict_svm)
##
## Data: predict_svm in 7486 controls (train_eval$Class 0) < 14 cases (train_eval$Class 1).
## Area under the curve: 0.9505
accuracy.meas(train_eval$Class, predict_mlr1)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_mlr1)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.800
## recall: 0.286
## F: 0.211
accuracy.meas(train_eval$Class, predict_poisson1)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_poisson1)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.500
```

```

## recall: 0.143
## F: 0.111
accuracy.meas(train_eval$Class, predict_logit1)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_logit1)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.778
## recall: 0.500
## F: 0.304
accuracy.meas(train_eval$Class, predict_backward)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_backward)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.800
## recall: 0.286
## F: 0.211
accuracy.meas(train_eval$Class, predict_forward)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_forward)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.800
## recall: 0.286
## F: 0.211
accuracy.meas(train_eval$Class, predict_decision)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_decision)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: 0.002
## recall: 1.000
## F: 0.002
accuracy.meas(train_eval$Class, predict_rforest)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_rforest)
##
## Examples are labelled as positive when predicted is greater than 0.5

```

```
##
## precision: 0.889
## recall: 0.571
## F: 0.348

accuracy.meas(train_eval$Class, predict_svm)

##
## Call:
## accuracy.meas(response = train_eval$Class, predicted = predict_svm)
##
## Examples are labelled as positive when predicted is greater than 0.5
##
## precision: NaN
## recall: 0.000
## F: NaN

varImpPlot(rforest)
```

